

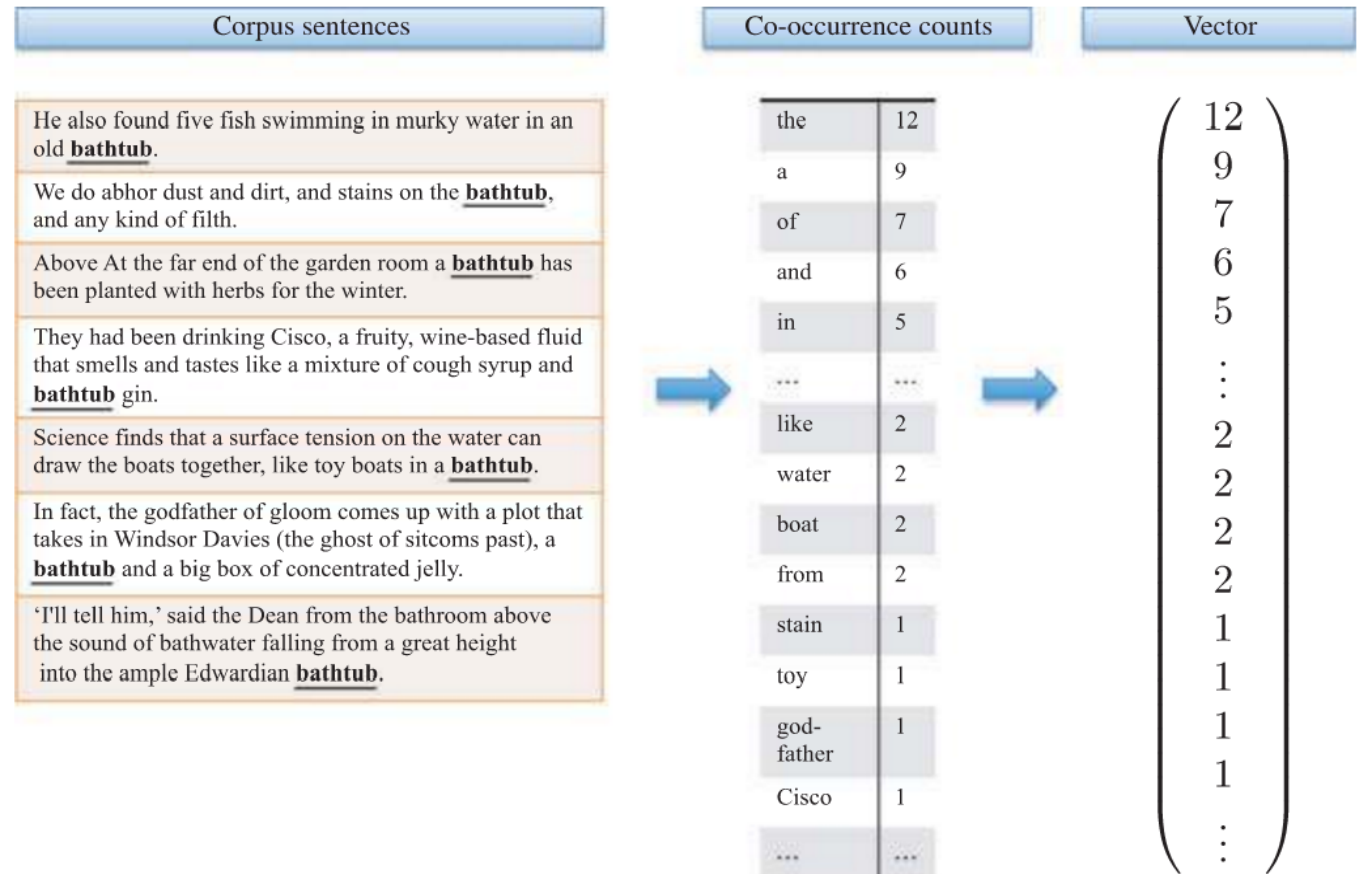
Lecture 6: Word Embeddings

The Distributional Hypothesis

- **The distributional hypothesis:** words that are used and occur in the same contexts tend to have similar meanings (Harris, 1954)
- Distributional semantics represents the meaning of words as a distribution over the word's contexts

Word Embeddings: Count-based Models

- Contexts are defined as neighboring words
 - Usually in a window of $\pm K$ words
- Dimensions correspond to context wordforms
- Values in entries correspond to counts – the number of times a word and a context word co-occurred




How Do These Vectors Represent Meaning?

- Similarity can be measured using vector distance metrics
- A popular choice is the “cosine similarity”:

$$\text{similarity}(w, u) = \frac{w \cdot u}{\|w\| \|u\|} = \frac{\sum_{i=1}^n w_i u_i}{\sqrt{\sum_{i=1}^n w_i^2} \sqrt{\sum_{i=1}^n u_i^2}}$$

How Do These Vectors Represent Meaning?

- Instead of representing words with 1-hot vectors (words are either the same or unrelated), embed them in a space that reflects their similarity patterns
- But taking neighboring wordforms as features is a bit naïve...



From Word Counts to Dimensionality Reduction

- Distributional semantics makes intuitive sense if we think of the dimensions as representing semantic features
- For example, a *dog* is a mammal, which is terrestrial, a carnivore and often domesticated
 - A *cat* is then more similar to a dog than a goat is, since they share these traits, while goats share only some
- However, neighboring words are considerably less abstract
 - For example, *car* and *automobile* are synonyms; but are represented as distinct dimensions
 - This fails to capture similarity between a word with *car* as a neighbor and a word with *automobile* as a neighbor

From Word Counts to Dimensionality Reduction

- One way to overcome this would be using dimensionality reduction methods
 - Such as singular value decomposition (Schütze, 1993) or the information bottleneck (Pereira et al., 1993)
- These are strong methods, used (with some variation) today as well
- However, we will devote the rest of the chapter to the more recent, prediction-based models

“Part of Speech Induction from Scratch”, Schütze, 1993

“Distributional clustering of English words”, Pereira, Tishby and Lee, 1993

Prediction-based Models

- Idea: instead of directly representing the distribution of a word, we can represent words as a vector from which the distribution of a word can be “decoded”
- **Task:** learn a network to predict a neighboring word from a given word
 - Sometimes called “self-supervision”
- An influential suite of methods for prediction-based embeddings is *word2vec* (Mikolov et al., 2013)
- We will review the basic implementation of the skip-gram model

Skip-gram (Setup)

- Notation:
 - Denote the j -th wordform in the vocabulary with x_j
 - Denote the vocabulary size (number of distinct wordforms) with V
 - N is a hyperparameter that determines the dimensionality of the embeddings


Skip-gram (Model)

The probability of predicting the neighbor x_j given the word x_k

Soft-max over the network's output

$$P(x_j|x_k) = \frac{e^{w'[:,j] \cdot w[k,:]}}{\sum_{j'=1,\dots,V} e^{w'[:,j'] \cdot w[k,:]}}$$

W and W' are the parameters of the model



Skip-gram (Training)

- Training is carried out by maximizing

$$\operatorname{argmax}_{W, W'} \log [P(\textit{text})] =$$

$$\operatorname{argmax}_{W, W'} \sum_{(x_k, x_j) \in \textit{text}} \log [P(x_j | x_k)]$$

where x_k and x_j are any pair of words no more than K tokens apart

Skip-gram


Each wordform now has two embeddings:

input embedding in the input matrix W

- Row j of the input matrix W is the N dimensional embedding for word j in the vocabulary.

output embedding v' , in output matrix W'


- Column j of the output matrix W' is a N dimensional vector embedding for word j in the vocabulary.
- The input and output embeddings are often concatenated to yield the word embedding for j



Word Embeddings as Features

- Word embeddings are often used as features to supervised learning tasks
- For instance, in text classification:

Instead of bag-of-words, one might input the sum of the embeddings of the words in the document



Word Embeddings as Features

- In this case, we don't really care what the dimensions represent
 - They are just useful feature representations, where the “semantics” of these features are unknown
- In fact, word embeddings are often used as initialization to the first layer of a neural network, and are later updated during training
 - This is sometimes called “fine-tuning”
- Applications cover all aspects of NLP. Essentially any work in NLP that use neural networks (and not only them), will use distributional embeddings to represent the words
 - We will see an example next lesson, and more towards the end of the course