





Sentiment Analysis and Introduction to RNNs

Partially based on slides by Jurafsky and Martin
Speech and Language Processing, 3rd Edition

Sentiment Analysis

Example #1: Movie Reviews

-  • Unbelievably disappointing
-  • Full of zany characters and richly applied satire, and some great plot twists
-  • This is the greatest screwball comedy ever filmed
-  • It was pathetic. The worst part about it was the boxing scenes.

Sentiment Analysis

Example #2: Product Reviews



HP Officejet 6500A Plus e-All-in-One Color Ink-jet - Fax / copier / printer / scanner

\$89 online, \$100 nearby ★★★★★ 377 reviews

September 2010 - Printer - HP - Inkjet - Office - Copier - Color - Scanner - Fax - 250 sh

Reviews

Summary - Based on 377 reviews




What people are saying

ease of use	<div><div></div><div></div></div>	"This was very easy to setup to four computers."
value	<div><div></div><div></div></div>	"Appreciate good quality at a fair price."
setup	<div><div></div><div></div></div>	"Overall pretty easy setup."
customer service	<div><div></div><div></div></div>	"I DO like honest tech support people."
size	<div><div></div><div></div></div>	"Pretty Paper weight."
mode	<div><div></div><div></div></div>	"Photos were fair on the high quality mode."
colors	<div><div></div><div></div></div>	"Full color prints came out with great quality."

- A Sentiment Analysis system called CALM applied to Twitter predicts the Dow Jones Industrial Average (DJIA) 3 days later
- Such algorithms are already used by hedge funds

Johan Bollen, Huina Mao, Xiaojun Zeng. 2011. [Twitter mood predicts the stock market](#), Journal of Computational Science 2:1, 1-8. 10.1016/j.jocs.2010.12.007.



Scherer Typology of Affective States

- **Emotion:** brief organically synchronized ... evaluation of a major event
 - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
 - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
 - *liking, loving, hating, valuing, desiring*
- **Personality traits:** stable personality dispositions and typical behavior tendencies
 - *nervous, anxious, reckless, morose, hostile, jealous*

Sentiment Analysis: Definition

- Sentiment analysis is the detection of **attitudes**
“enduring, affectively colored beliefs, dispositions towards objects or persons”
 1. **Holder (source)** of attitude
 2. **Target (aspect)** of attitude
 3. **Type** of attitude
 - From a set of types: *like, love, hate, value, desire*, etc.
 - Or (more commonly) simple weighted **polarity**: *positive, negative, neutral*, together with *strength*
 4. **Text** containing the attitude
 - Sentence or entire document

Sentiment Analysis

- Simplest task:
 - Is the attitude of this text positive or negative?
- More complex:
 - Rank the attitude of this text from 1 to 5
- Advanced:
 - Detect the target, source, or complex attitude types
- Sometimes not only complete sentences, but also their sub-parts receive a sentiment value
 - [The pizza was overall bad], but the toppings were fun! → overall negative

Sentiment Classification in Movie Reviews

- Is an IMDB movie review positive or negative?



when _star wars_ came out some twenty years ago
, the image of traveling throughout the stars has
become a commonplace image . [...]

when han solo goes light speed , the stars change
to bright lines , going towards the viewer in lines
that converge at an invisible point .

cool .



“ snake eyes ” is the most aggravating
kind of movie : the kind that shows so
much potential then becomes
unbelievably disappointing .

it's not just because this is a brian
depalma film , and since he's a great
director and one who's films are always
greeted with at least some fanfare .

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79—86.

Bo Pang and Lillian Lee. 2004. A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts. ACL, 271-278

A Simple Classifier

- Log-linear or Naïve Bayes classifier
- Features:
 - Tokenized words
 - Possibly mark-up (e.g., hashtags in Twitter, headers in HTMLs)
- Features are often binary
 - Indicating whether a word appeared or did not appear in the document (bag of words)
 - Often works better for text classification than word frequency
 - The effect of frequency is non-linear

Error Analysis:

What makes reviews hard to classify?

- **Sarcasm:**

- Perfume review in *Perfumes: the Guide*:
 - “If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut.”
- On *Automobile Steering Wheel Attachable Work Surface*:
 - “You wouldn’t believe how much more interesting my commute is now that I have something to do other than just stare out the window! I’m using it right now to post this review and I never”



- **Thwarted Expectations and Ordering Effects:**

- “This film should be brilliant. It sounds like a great plot, the actors are first grade, and the supporting cast is good as well, and Stallone is attempting to deliver a good performance. However, it can’t hold up.”

Negation in Sentiment Analysis

- One practice is to add NOT_ to every word between negation and following punctuation:

didn't like this movie , but I



didn't NOT_like NOT_this NOT_movie but I

Negation in Sentiment Analysis

- This is a very crude solution:
 - Explicit negation is only one way to reverse polarity
 - “He **failed** to convey the importance of his message”
 - Negation scope:
 - “I didn’t like the exposition in this otherwise excellent film”
 - Logical structure can be more complex
 - “I don’t think anyone could have done a better job”
- More recent approaches take the context (surrounding words) into account

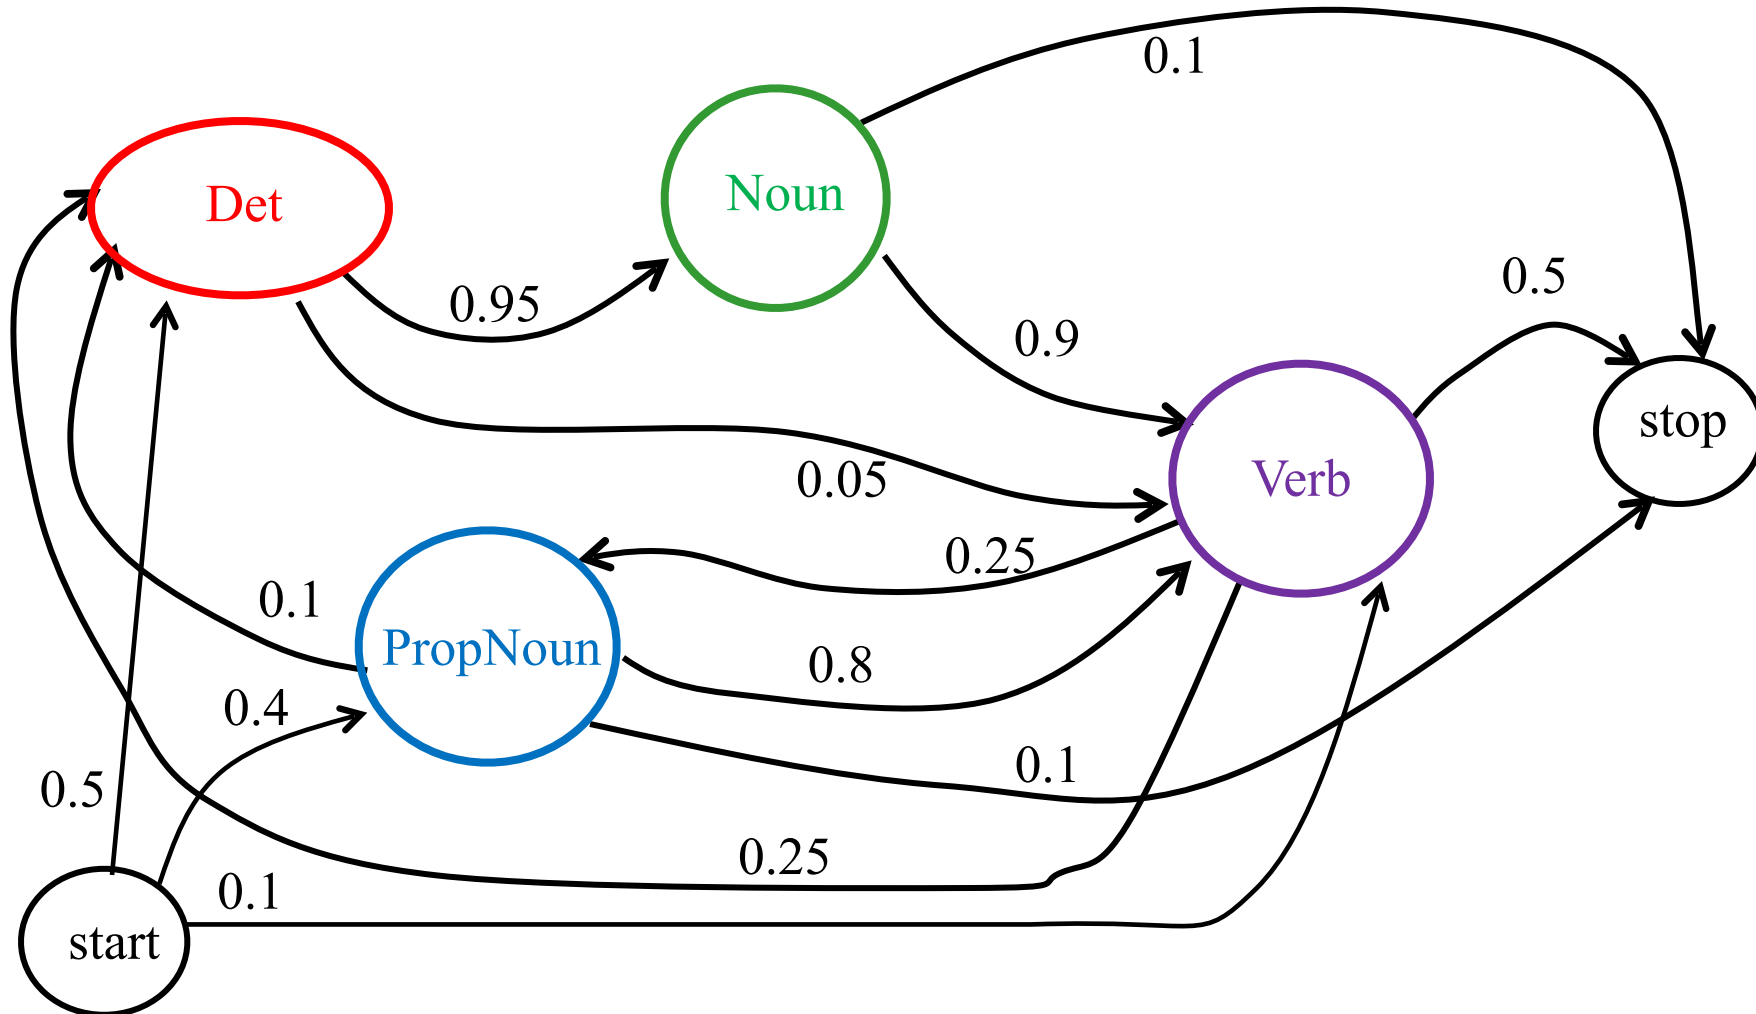
Sentiment Analysis as Sequence Labeling

- This construal of sentiment analysis attempts to capture the meaning of a word in context by encoding (parts of) the sentence as features
- Recently: using Recurrent Neural Networks (RNNs)
- Recall the underlying assumption in Markov (n-gram) models:
 - It's enough to know the last n tokens you've encountered to know what's next
 - Alternative view: the probability of a sequence is the product of the probabilities of its *n-grams*

Sentiment Analysis as Sequence Labeling

- Consider the example:
 - *How can you not see this movie?*
 - *You should not see this movie*
 - Great idea, you have a real product development team!
 - Great idea, now try again with a real product development team!
- How well will a bi-gram model work?
 - Similar unigrams and bigrams → similar prediction
- The problem with Markov Models: need to maintain a **state** to capture distant influences
 - The size of the space increases exponentially with the order of the model

Recall: Markov Models are Finite State Automata (FSAs) with transition probabilities



Recurrent Neural Networks

- Motivation:
 - Neural network model, but with a state
 - How can we borrow ideas from FSAs?
- RNNs are FSAs ...
 - With a twist
 - No longer finite in the same sense
 - The state is an embedding of the history in a continuous space
 - The embedding function of the history to a vector is learned as well

Recurrent Neural Networks

- Map from dense sequence to dense representation
 - Maps a sequence of vector inputs to a sequence of vector states

$$x_1, \dots, x_n \rightarrow s_1, \dots, s_n$$

- A (parametrized) transition function R does the mapping:

$$s_i = R(s_{i-1}, x_i)$$

- R is parametrized and parameters are shared across steps

$$s_4 = R(s_3, x_4) = R(R(s_2, x_3), x_4) = R(R(R(s_1, x_2), x_3), x_4)$$

Recurrent Neural Networks

- An output function O maps states to (vector) outputs, which are often viewed as a distribution over the possible labels

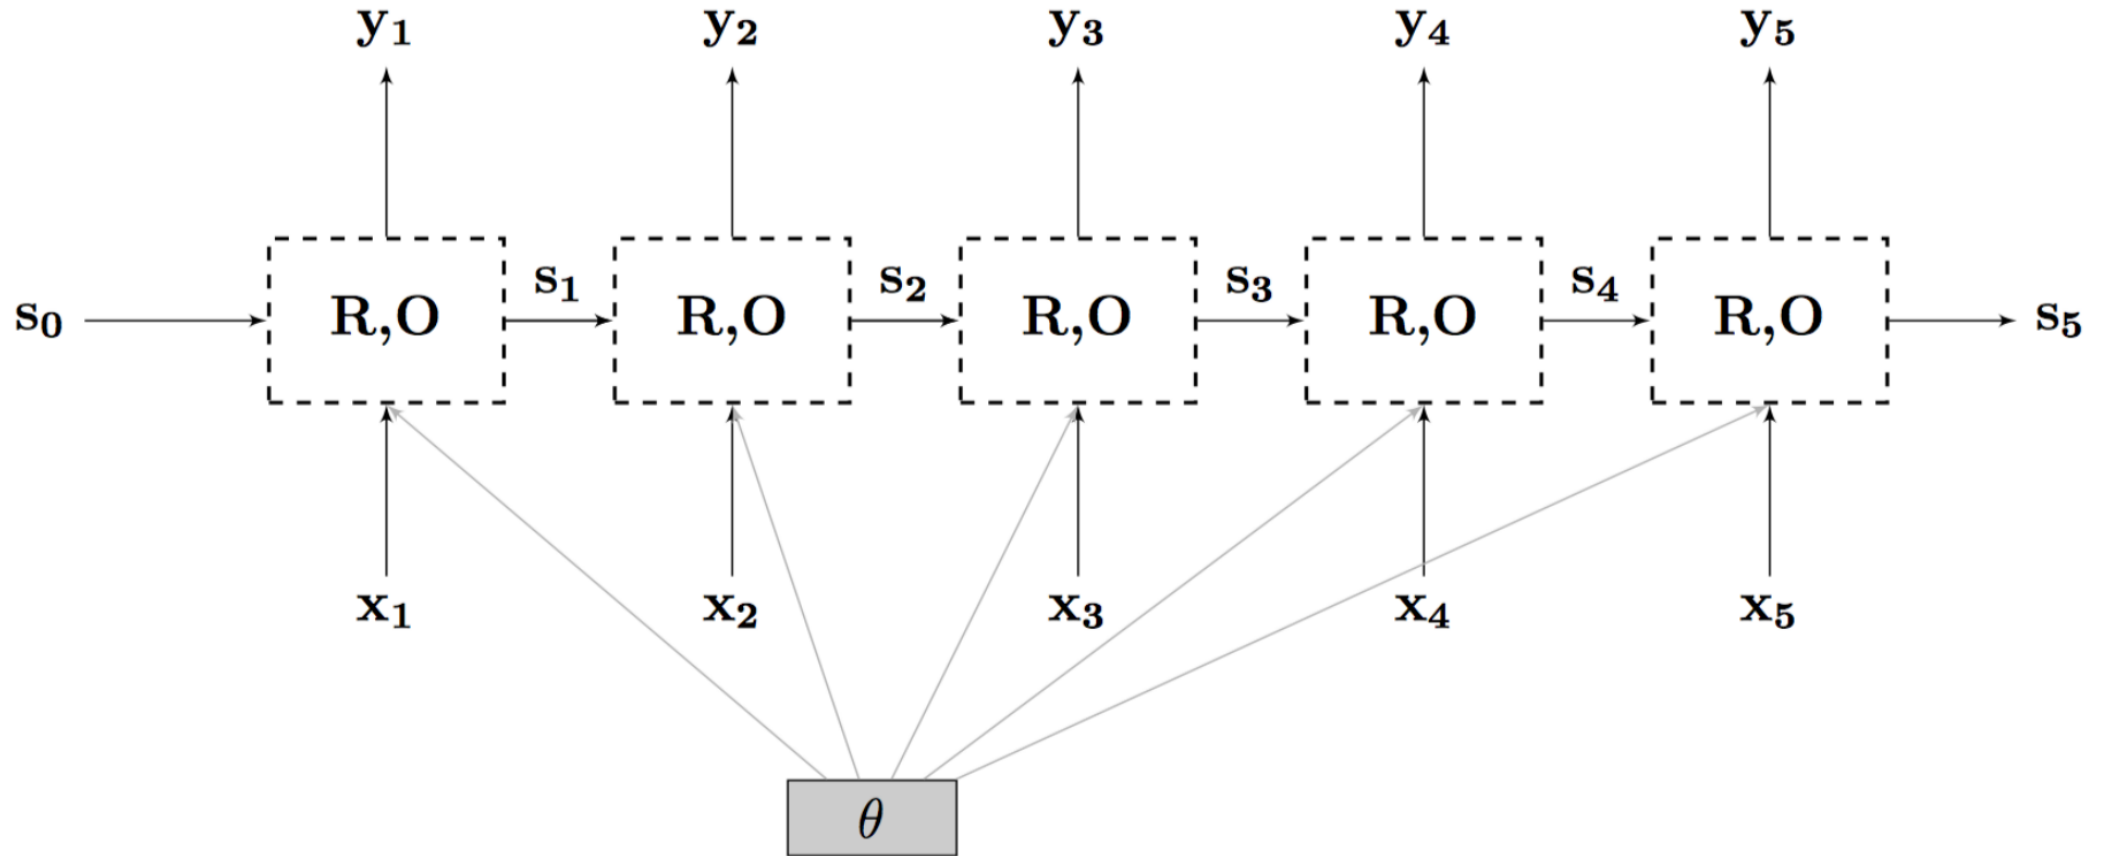
- Example:

$$R_{Elman}(s, x) = \tanh(W[x, s] + b)$$

$$O(s_i) = \text{softmax}(W's_i + b')$$

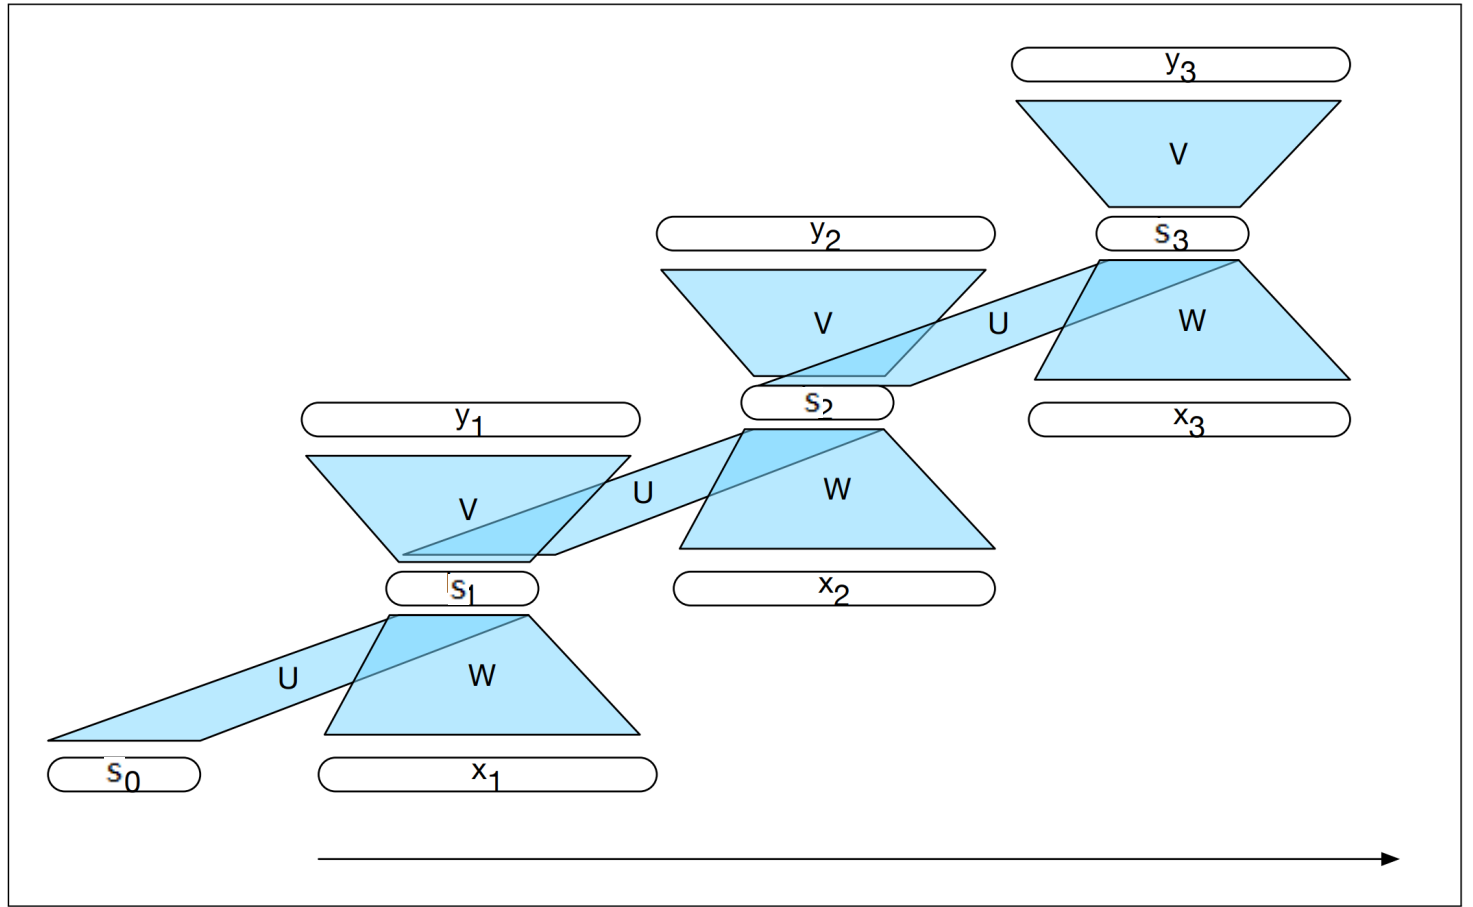
- Conceptually:
 - R is the function embedding the history in a vector
 - O is the function relating the embedded history and the output
 - They are learned jointly

RNNs: Graphical Representation



Recurrent Neural Language Models

- An “unrolled” view of an RNN
- The network stays the same at each time stamp. What’s different are the inputs and previous hidden states




Back to Sentiment Analysis

- When using RNNs for sentence classification (such as sentiment analysis), it is often practical to use Bi-RNNs
- Bi-RNNs:
 - 2 RNNs, one going back to forth and the other forth to back
 - Output function is a function of both states
- This allows the states associated with each word to encode relevant information from the words following them and preceding them

$$\vec{h}_t = (W_{x\vec{h}}x_t + W_{\vec{h}\vec{h}}\vec{h}_{t-1} + b_{\vec{h}}) \quad (9)$$

$$\overleftarrow{h}_t = (W_{x\overleftarrow{h}}x_t + W_{\overleftarrow{h}\overleftarrow{h}}\overleftarrow{h}_{t+1} + b_{\overleftarrow{h}}) \quad (10)$$

$$y_t = W_{\vec{h}y}\vec{h}_t + W_{\overleftarrow{h}y}\overleftarrow{h}_t + b_y \quad (11)$$



Back to Sentiment Analysis

- One simple way to do sentiment analysis (or other sentence classification) with Bi-RNNs is to average the output sequence:

$$y = \frac{1}{N} \sum_i y_i$$


- Now train a binary (log-linear) classifier for predicting the sentiment:

$$P(+|y) = \frac{1}{1+e^{-(w^t \cdot y + b)}} = \text{sigmoid}(w^t \cdot y + b)$$

$$\vec{h}_t = (W_{x\vec{h}}x_t + W_{\vec{h}\vec{h}}\vec{h}_{t-1} + b_{\vec{h}}) \quad (9)$$

$$\overleftarrow{h}_t = (W_{x\overleftarrow{h}}x_t + W_{\overleftarrow{h}\overleftarrow{h}}\overleftarrow{h}_{t+1} + b_{\overleftarrow{h}}) \quad (10)$$

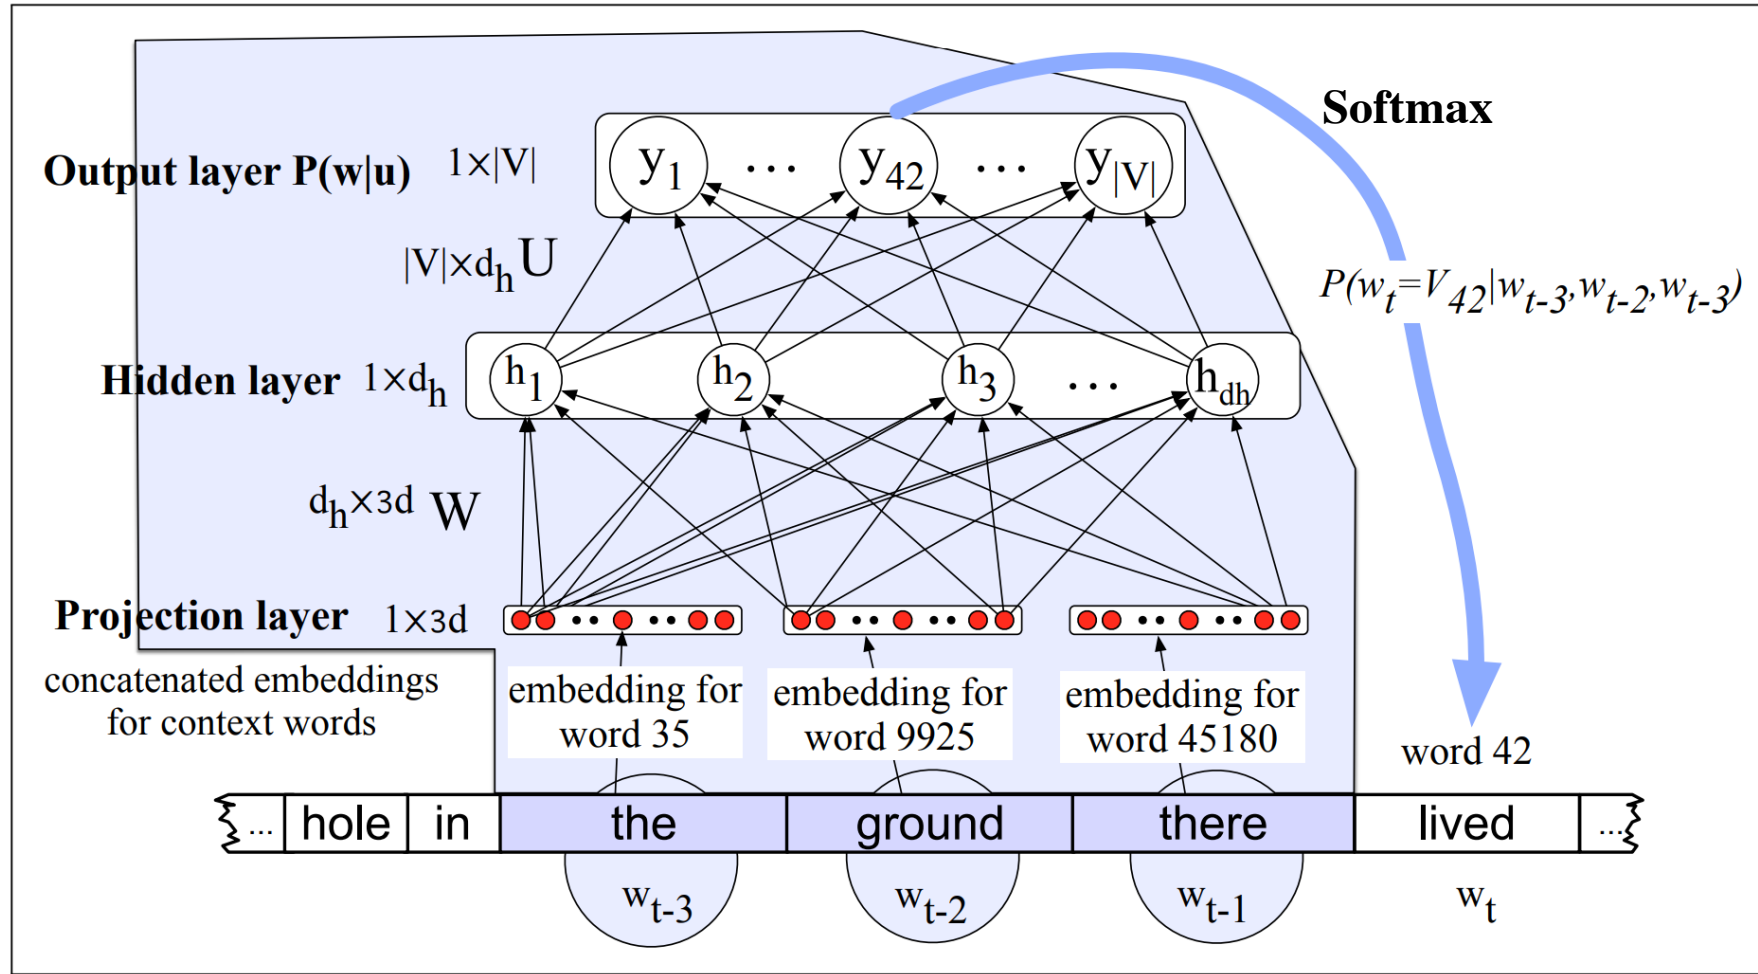
$$y_t = W_{\vec{h}y}\vec{h}_t + W_{\overleftarrow{h}y}\overleftarrow{h}_t + b_y \quad (11)$$



Context in RNN models for Sentiment Analysis

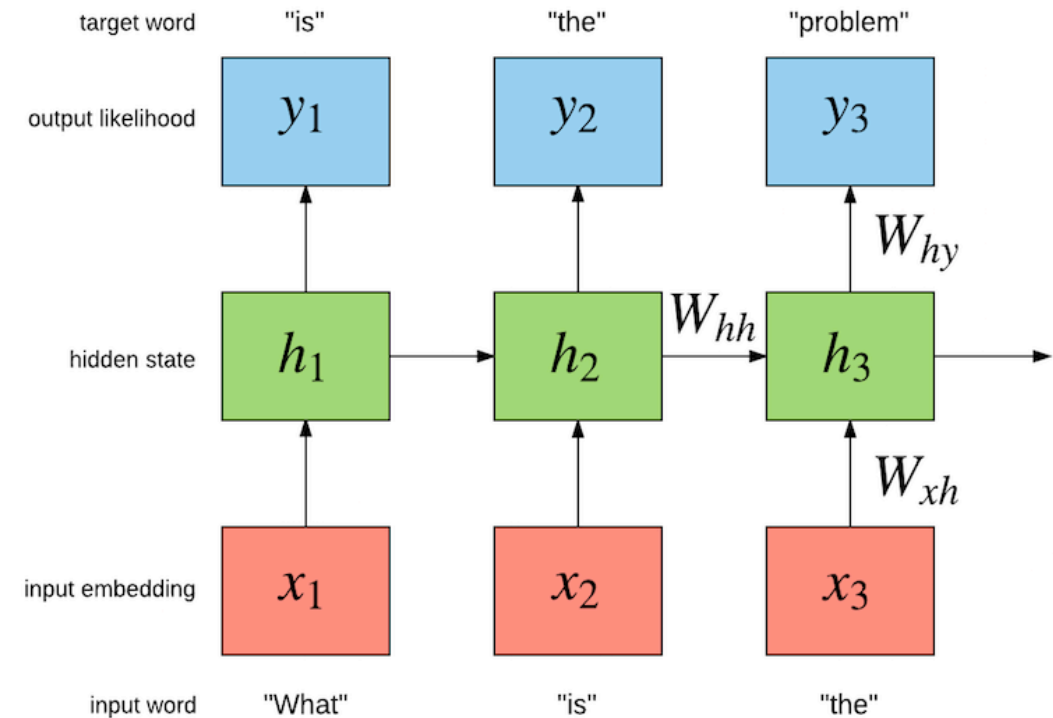
- RNN models are able to take the context (both preceding and following) into account, as well as the linear order between the words
 - bag-of-words models cannot
- They indeed show better performance in tasks such as sentiment analysis
- However, the state sequence is not easy to interpret
 - Much heavier computationally, especially training
- Further investigation is needed to establish what contextual and semantic aspects of sentences are captured using these techniques

Feed Forward Neural Language Models



Recurrent Neural Language Models

- Language models based on RNNs have shown much power in recent years, consistently surpassing n-gram models
- The basic architecture is that of a sequence recurrent neural net (RNN)
- Just like in Feed Forward LMs:
 - Input words are converted embeddings
 - The output is passed through a softmax layer, which defines the probability of the next word
- The difference is that the hidden state is passed as input to the new layer



Reminder: Loss Functions

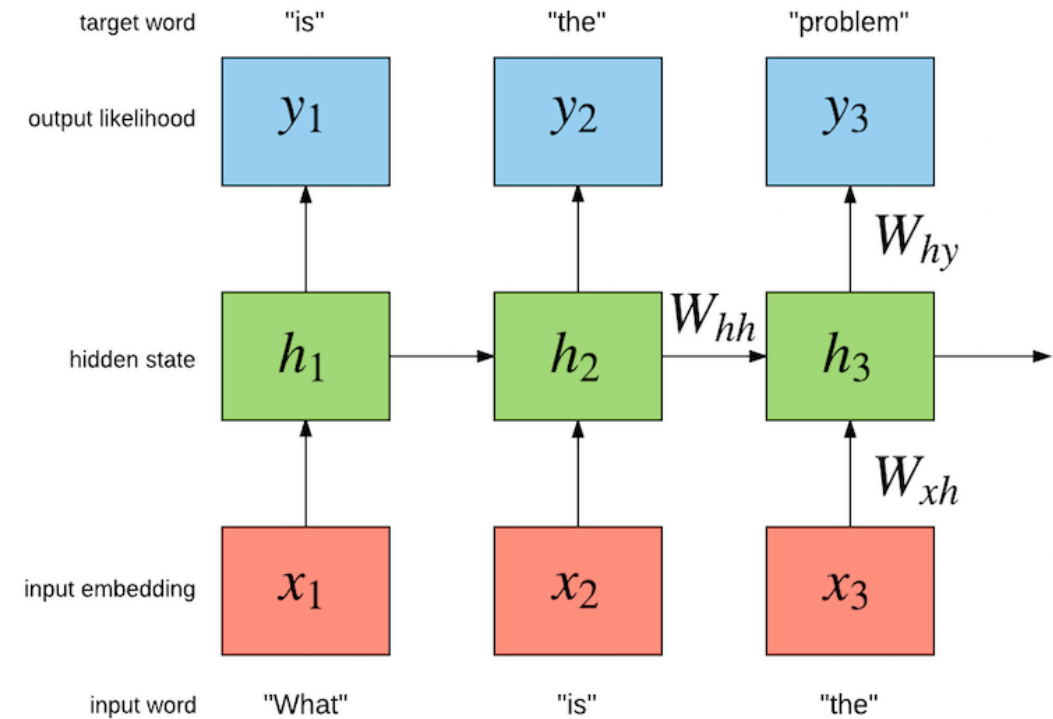
- Training is done by defining a loss function (i.e. a function that determines how “bad” a classification is relative to the gold standard)
- The network then attempts to minimize the empirical risk. For a sample $S=\{x_1, \dots, x_m\}$ with gold standard labels $\{y_1, \dots, y_m\}$, the empirical risk is:

$$L_S(\theta) = \frac{1}{m} \sum_{i=1}^m \ell(\theta; (x_i, y_i))$$

Recurrent Neural Language Models

- Neural language models are generally trained to maximize the sum of the predicted log-probabilities of the correct words
→ the loss function is the negative predicted log-probabilities of the correct words
- In the example to the right, the loss will be

$$-\log(\text{soft_max}(y_1)(is)) - \log(\text{soft_max}(y_2)(the)) - \log(\text{soft_max}(y_3)(problem))$$



Character-level Language Models

- Vocabulary: characters instead of words
- Advantage:
 - Small vocabulary → compact model
 - Can generalize over morphologically similar words
- However:
 - Need to learn how to spell
 - Longer range dependencies between tokens

Character-level Language Models

